

Machine Learning Exercises - IV - Solutions

May 11, 2022

Kannan Singaravelu

* * *

Exercise 7

Perform cluster analysis on DJIA Index components using K-Means and build an equal weight portfolio from the selected stocks. Retrieve the list of DJIA components; identify and clean any missing data points. Cluster stocks based on weekly ATR and compare it with the original dataset.

Solutions

K-Means Clustering

Clustering is a branch of unsupervised machine learning models that seeks to learn from the properties of the data by identifying groups or clusters in the dataset.

The k -means algorithm searches for a predetermined number of clusters within an unlabeled dataset and is based on the assumptions that the optimal cluster will have cluster center and each point is closer to its own cluster center than to other cluster centers.

```
[ ]: # Ignore warnings
import warnings
warnings.filterwarnings('ignore')

# Import Libraries
import pandas as pd
import numpy as np
import pyfolio as pf

from kneed import KneeLocator
import matplotlib.pyplot as plt

from sklearn.cluster import KMeans
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MinMaxScaler

[ ]: # Load the pre-saved data dict
ohlcv = np.load('dow_ohlcv.npy', allow_pickle='TRUE').item()
ohlcv['MMM'].head()
```

```
[ ]:
      Open      High      Low      Close      Volume
Date
2009-12-31  62.906018  63.123738  61.967570  62.065166  2049800
2010-01-04  62.380479  62.650753  62.065162  62.327927  3043700
2010-01-05  62.162779  62.485606  61.336941  61.937550  2847000
2010-01-06  62.973588  63.514135  62.695811  62.815929  5268500
2010-01-07  62.553137  62.883472  61.652230  62.860950  4470100
```

```
[ ]: # List of DJIA stocks
dow_stocks = ['MMM', 'AXP', 'AMGN', 'AAPL', 'BA', 'CAT', 'CVX', 'CSCO', 'KO',
              ↪ 'DOW', 'GS', 'HD', 'HON', 'IBM', 'INTC', 'JNJ', 'JPM', 'MCD', 'MRK', 'MSFT',
              ↪ 'NKE', 'PG', 'CRM', 'TRV', 'UNH', 'VZ', 'V', 'WBA', 'WMT', 'DIS']
```

```
[ ]: # Function to calculate average true range
def ATR(df,n):
    "function to calculate Average True Range"
    df = df.copy()

    df['H-L'] = abs(df['High']-df['Low'])
    df['H-PC'] = abs(df['High']-df['Close'].shift(1))
    df['L-PC'] = abs(df['Low']-df['Close'].shift(1))

    df['TR'] = df[['H-L', 'H-PC', 'L-PC']].max(axis=1,skipna=False)
    df['ATR'] = df['TR'].rolling(n).mean()

    df2 = df.drop(['H-L', 'H-PC', 'L-PC'],axis=1)
    return df2['ATR']
```

```
[ ]: # Add ATR for each stocks
for symbol in dow_stocks:
    ohlc[symbol]['ATR'] = ATR(ohlc[symbol],21)
```

```
[ ]: # Subsume into dataframe
df = pd.DataFrame({symbol: ohlc[symbol]['ATR'] for symbol in dow_stocks})

# Check for missing values
df.isnull().sum()
```

```
[ ]: MMM      22
     AXP      22
     AMGN     22
     AAPL     22
     BA       22
     CAT      43
     CVX      22
     CSCO     22
     KO       22
```

```

DOW      2340
GS        22
HD        22
HON       22
IBM       22
INTC      22
JNJ       22
JPM       22
MCD       22
MRK       22
MSFT      22
NKE       22
PG        22
CRM       22
TRV       22
UNH       22
VZ        22
V         22
WBA       22
WMT       22
DIS       22
dtype: int64

```

```

[ ]: # Fill forward the missing values and drop DOW company from the list
df.fillna(method='bfill', axis=0, inplace=True)
df.drop(['DOW'], axis=1, inplace=True)

```

```

[ ]: # Resample to a weekly timeframe for cluster analysis
px = df.resample('W-FRI').mean()
px = px.T

# Check output
px.head(2)

```

```

[ ]: Date    2010-01-01    2010-01-08    2010-01-15    2010-01-22    2010-01-29    2010-02-05    \
MMM         1.113625     1.113625     1.113625     1.113625     1.113625     1.153880
AXP         1.035185     1.035185     1.035185     1.035185     1.035185     1.046649

Date    2010-02-12    2010-02-19    2010-02-26    2010-03-05    ...    2020-07-31    \
MMM         1.233820     1.308653     1.232967     1.071002    ...         3.304308
AXP         1.079262     1.101860     0.908801     0.774188    ...         2.843043

Date    2020-08-07    2020-08-14    2020-08-21    2020-08-28    2020-09-04    2020-09-11    \
MMM         3.337810     3.339979     3.176965     2.863492     2.869821     3.166674
AXP         2.639048     2.661429     2.715334     2.638858     2.682953     2.666072

Date    2020-09-18    2020-09-25    2020-10-02

```

MMM	3.413893	4.169714	4.322380
AXP	2.634953	3.009524	3.067619

[2 rows x 562 columns]

Elbow Plot The number of clusters is a hyperparameter to clustering models and choose the optimal number of clusters is critical for the model. We identify the elbow point programmatically for this exercise.

```
[ ]: scaler = MinMaxScaler()
scaled_px = scaler.fit_transform(px)
```

```
[ ]: # Get the inertia
sse = []
for k in range(1, 30):
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(scaled_px)
    sse.append(kmeans.inertia_)

# Knee Locator
kl = KneeLocator(range(1, 30), sse, curve="convex", direction="decreasing")
kl.elbow
```

```
[ ]: 6
```

Build Clusters

We will now fit the cluster model.

```
[ ]: # Build clusters
model = KMeans(n_clusters=6)
model.fit(scaled_px)

labels = model.predict(scaled_px)
labels
```

```
[ ]: array([4, 1, 2, 3, 5, 4, 4, 3, 3, 0, 1, 1, 4, 3, 1, 1, 1, 3, 3, 3, 3, 4,
        1, 5, 3, 1, 1, 3, 1])
```

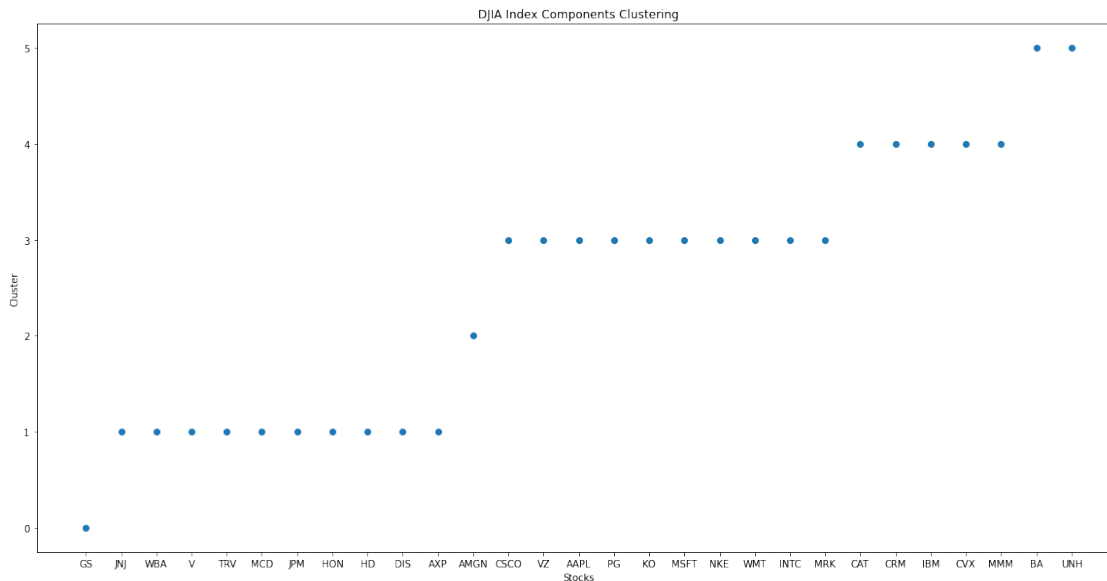
```
[ ]: # Remove DOW from the list
companies = dow_stocks
companies.remove('DOW')
```

```
[ ]: df1 = pd.DataFrame({'Cluster': labels,
                        'Companies': companies,
                        'ATR': px.mean(axis=1),
                        }).sort_values(by=['Cluster'], axis = 0)
```

```
df1 = df1.reset_index(drop=True)
df1
```

```
[ ]:      Cluster Companies      ATR
0         0      GS  3.640206
1         1     JNJ  1.306309
2         1     WBA  1.119627
3         1       V  1.501398
4         1     TRV  1.423909
5         1     MCD  1.650644
6         1     JPM  1.310115
7         1     HON  1.606682
8         1      HD  1.987495
9         1     DIS  1.486680
10        1     AXP  1.424071
11        2    AMGN  2.664776
12        3    CSCO  0.538751
13        3      VZ  0.606684
14        3    AAPL  0.680922
15        3      PG  1.030430
16        3      KO  0.493262
17        3    MSFT  1.290491
18        3     NKE  0.969232
19        3     WMT  1.116245
20        3    INTC  0.695828
21        3     MRK  0.865679
22        4     CAT  2.079966
23        4     CRM  2.199421
24        4     IBM  2.095919
25        4     CVX  1.630457
26        4     MMM  2.101291
27        5      BA  3.791872
28        5     UNH  2.748506
```

```
[ ]: # Plot Clusters
plt.figure(figsize=(20,10))
plt.scatter(df1.Companies, df1.Cluster)
plt.xlabel('Stocks')
plt.ylabel('Cluster')
plt.title('DJIA Index Components Clustering');
```



Portfolio Construction Shortlisting the six stocks from the above clusters (one for each), we will now build a portfolio and compare the returns with all stock portfolio.

```
[ ]: # Cluster portfolio stocks
portfolio_stocks = ['GS', 'JNJ', 'UNH', 'CSCO', 'CRM', 'AMGN']
port = pd.DataFrame({symbol: ohlc[symbol]['Close'] for symbol in
    ↳ portfolio_stocks})
port.dropna(inplace=True)

port
```

```
[ ]:
```

	GS	JNJ	UNH	CSCO	CRM \
Date					
2009-12-31	144.296997	46.714996	25.778193	18.048286	18.442499
2010-01-04	147.920776	46.910812	26.666224	18.613705	18.705000
2010-01-05	150.535919	46.366856	26.623932	18.530785	18.625000
2010-01-06	148.929138	46.743999	26.886124	18.410154	18.592501
2010-01-07	151.843475	46.410374	27.917923	18.493084	18.510000
...
2020-09-23	186.119995	144.440002	292.140015	37.930145	235.990005
2020-09-24	195.110001	144.669998	292.660004	37.504074	237.550003
2020-09-25	194.949997	145.660004	302.500000	38.098591	242.740005
2020-09-28	199.070007	147.110001	303.230011	38.772377	246.669998
2020-09-29	196.789993	147.059998	304.149994	38.703018	247.449997

```

AMGN
Date
```

```

2009-12-31    45.404182
2010-01-04    46.327202
2010-01-05    45.925884
2010-01-06    45.580757
2010-01-07    45.163395
...
2020-09-23    242.589996
2020-09-24    240.320007
2020-09-25    243.820007
2020-09-28    247.029999
2020-09-29    248.300003

```

[2705 rows x 6 columns]

```

[ ]: # Calculate portfolio returns
portfolio_returns = port.pct_change().fillna(0)
port['Returns'] = portfolio_returns.mean(axis=1)
port.head(2)

```

```

[ ]:
      GS      JNJ      UNH      CSC0      CRM      AMGN  \
Date
2009-12-31  144.296997  46.714996  25.778193  18.048286  18.442499  45.404182
2010-01-04  147.920776  46.910812  26.666224  18.613705  18.705000  46.327202

      Returns
Date
2009-12-31  0.000000
2010-01-04  0.021607




```

```

[ ]: # Create Tear sheet using pyfolio
pf.create_simple_tear_sheet(port['Returns'])

```

<IPython.core.display.HTML object>




```
[ ]: # All stocks portfolio
all_stocks = pd.DataFrame({symbol: ohlc[symbol]['Close'] for symbol in
    ↪companies})
all_stocks.dropna(inplace=True)
all_stocks.head(2)
```

```
[ ]:
      MMM      AXP      AMGN      AAPL      BA      CAT \
Date
2009-12-31  62.065166  34.434761  45.404182  6.503574  41.856789  41.498463
2010-01-04  62.327927  34.774685  46.327202  6.604801  43.441975  42.634403

      CVX      CSC0      KO      GS ...      NKE \
Date
2009-12-31  50.924435  18.048286  18.951757  144.296997  ...  12.066024
2010-01-04  52.293617  18.613705  18.965061  147.920776  ...  11.934528

      PG      CRM      TRV      UNH      VZ      V \
Date
2009-12-31  43.431492  18.442499  38.205708  25.778193  18.633041  17.935587
2010-01-04  43.782478  18.705000  38.167397  26.666224  18.717409  18.075037

      WBA      WMT      DIS
Date
2009-12-31  28.350834  40.954620  28.090706
2010-01-04  28.798639  41.552284  27.933924

[2 rows x 29 columns]
```

```
[ ]: # Calculate all stocks portfolio returns
all_stocks_returns = all_stocks.pct_change().fillna(0)
all_stocks['Returns'] = all_stocks_returns.mean(axis=1)
all_stocks.head(2)
```

```
[ ]:
      MMM      AXP      AMGN      AAPL      BA      CAT \
Date
2009-12-31  62.065166  34.434761  45.404182  6.503574  41.856789  41.498463
2010-01-04  62.327927  34.774685  46.327202  6.604801  43.441975  42.634403

      CVX      CSC0      KO      GS ...      PG \
Date
2009-12-31  50.924435  18.048286  18.951757  144.296997  ...  43.431492
2010-01-04  52.293617  18.613705  18.965061  147.920776  ...  43.782478




      CRM      TRV      UNH      VZ      V      WBA \
Date
2009-12-31  18.442499  38.205708  25.778193  18.633041  17.935587  28.350834
2010-01-04  18.705000  38.167397  26.666224  18.717409  18.075037  28.798639
```

	WMT	DIS	Returns
Date			
2009-12-31	40.954620	28.090706	0.000000
2010-01-04	41.552284	27.933924	0.013946

[2 rows x 30 columns]

```
[ ]: # Create Tear sheet using pyfolio
pf.create_simple_tear_sheet(all_stocks['Returns'])
plt.show()
```

<IPython.core.display.HTML object>



Conclusion The cluster stocks generated a CAGR of 17% with a maximum drawdown of ~30% while the all stocks portfolio consisting of 29 DJIA index stocks generated an annualized return of ~16% with a maximum drawdown of 33%. This study highlight that with 20% of all stocks, we can construct a portfolio that can outperform the all stocks portfolio with an alpha of 59% and an improved sortino ratio.

Note:

1. The data is not treated for in/out sample as the objective here is to showcase the application of clustering methods. Accordingly, the actual results may vary.
2. Arbitrary selection of Cluster stocks can be avoided by adopting a minimum distance measure in stock selection.

References

- [Scikit-learn K-Means Clustering](#)
- [Pyfolio-reloaded](#)
- [Python resources](#)